

Question 1

Since `biomarker-raw.csv` contains protein concentration levels across various samples, let's examine the distribution of a sample of these protein values.

```
library(ggplot2)
library(reshape2)

data <- read.csv("data2/biomarker-raw.csv")

# Convert the selected protein columns to numeric (skip non-numeric values)
proteins <- data[, 3:7] # Select a sample of 5 protein columns
proteins <- data.frame(lapply(proteins, function(x) as.numeric(as.character(x))))

## Warning in FUN(X[[i]], ...): NAs introduced by coercion
## Warning in FUN(X[[i]], ...): NAs introduced by coercion
## Warning in FUN(X[[i]], ...): NAs introduced by coercion
## Warning in FUN(X[[i]], ...): NAs introduced by coercion
## Warning in FUN(X[[i]], ...): NAs introduced by coercion

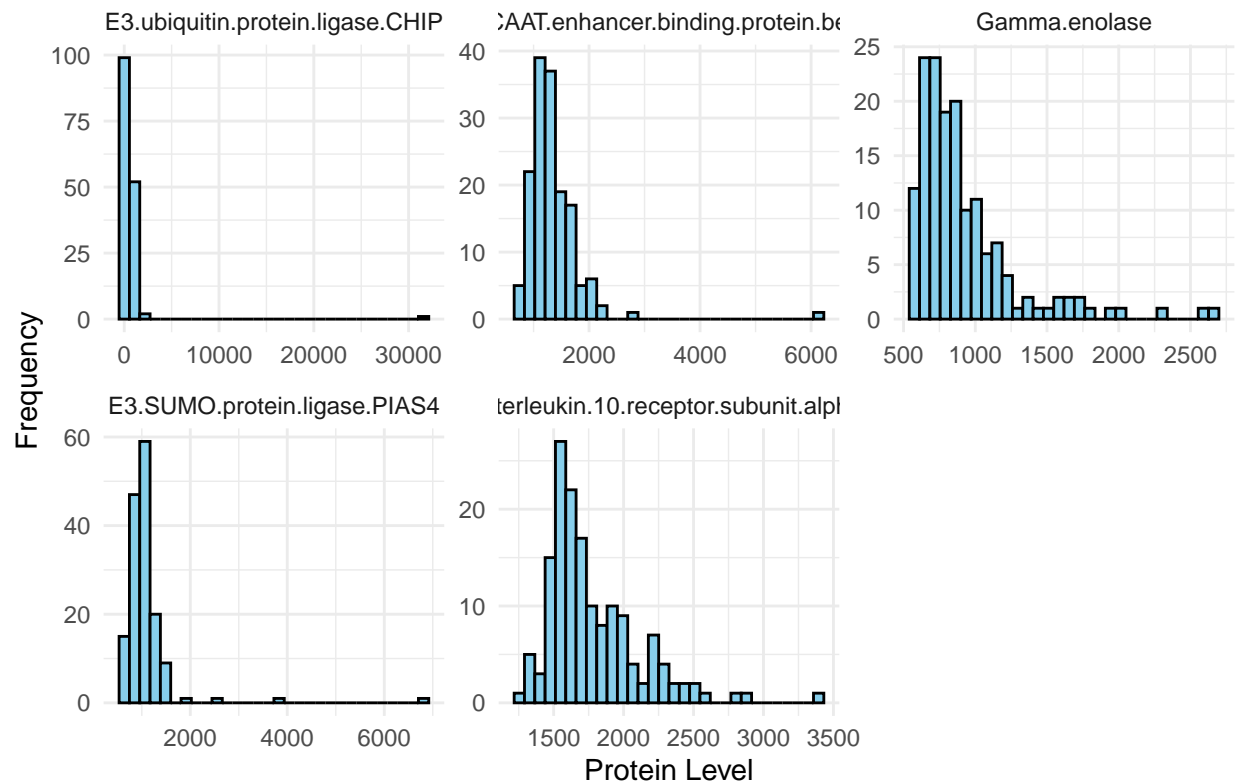
# Melt the data for ggplot
proteins_long <- melt(proteins, variable.name = "Protein", value.name = "Level")

## No id variables; using all as measure variables

# Plot histograms for each selected protein
ggplot(proteins_long, aes(x = Level)) +
  geom_histogram(bins = 30, color = "black", fill = "skyblue") +
  facet_wrap(~Protein, scales = "free") +
  labs(title = "Distribution of Raw Protein Levels (Sample Proteins)",
       x = "Protein Level", y = "Frequency") +
  theme_minimal()

## Warning: Removed 10 rows containing non-finite outside the scale range
## ('stat_bin()').
```

Distribution of Raw Protein Levels (Sample Proteins)



The histograms reveal that the distributions of raw protein levels are skewed right, with some values extending to higher ranges. This skewness is a common reason to apply a log transformation, which can help normalize these distributions.

Question 2

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(tidyr)
```

```
##
## Attaching package: 'tidyr'
```

```

## The following object is masked from 'package:reshape2':
##
## smiths

# Load necessary libraries
library(dplyr)
library(tidyr)

# Load the data (assuming it's processed after removing outlier trimming)
data <- read.csv("data2/biomarker-raw.csv")

# Convert relevant columns to numeric, skipping metadata columns
protein_data <- data[, -c(1,2)] # Adjust index based on the actual data structure

# Define a function to identify outliers based on  $\pm 3$  standard deviations
identify_outliers <- function(x) {
  upper_limit <- mean(x, na.rm = TRUE) + 3 * sd(x, na.rm = TRUE)
  lower_limit <- mean(x, na.rm = TRUE) - 3 * sd(x, na.rm = TRUE)
  return(x < lower_limit | x > upper_limit)
}

# Apply the outlier function across all protein columns
outliers <- protein_data %>%
  mutate(across(everything(), identify_outliers)) %>%
  rowwise() %>%
  mutate(outlier_count = sum(c_across(everything()), na.rm = TRUE)) %>%
  ungroup()

## Warning: There were 5268 warnings in 'mutate()'.
## The first warning was:
## i In argument: 'across(everything(), identify_outliers)'.
## Caused by warning in 'mean.default()':
## ! argument is not numeric or logical: returning NA
## i Run 'dplyr::last_dplyr_warnings()' to see the 5267 remaining warnings.

# Merge with subject ID and group columns
outlier_summary <- data %>%
  select(Group) %>%
  bind_cols(outliers["outlier_count"])

# Summarize the number of outlying values per subject and by group
outlier_by_group <- outlier_summary %>%
  group_by(Group) %>%
  summarise(avg_outliers = mean(outlier_count, na.rm = TRUE),
            total_outliers = sum(outlier_count, na.rm = TRUE),
            subjects_with_outliers = sum(outlier_count > 0))

# View results
print(outlier_by_group)

## # A tibble: 3 x 4
##   Group avg_outliers total_outliers subjects_with_outliers
##   <chr>      <dbl>          <int>          <int>

```

## 1 ""	0	0	0
## 2 "ASD"	0.0263	2	2
## 3 "TD"	0	0	0

Are there specific subjects (not values) that seem to be outliers? Yes, there are specific subjects with outliers. According to the table of nsummary, there are 2 subjects with outlying values in the ASD group.

Are outliers more frequent in one group or the other? Outliers are more frequent in the ASD group. ASD group with a total of 2 outliers, while the TD group has none.

Question 3

```
library(dplyr)
library(caret) #data partition
```

```
## Loading required package: lattice
```

```
library(dplyr) #ttest
library(purrr) #ttest
```

```
##
## Attaching package: 'purrr'
```

```
## The following object is masked from 'package:caret':
##
## lift
```

```
library(broom) #ttest
```

```
## Warning: package 'broom' was built under R version 4.4.1
```

```
library(tidyverse) #ttest
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v forcats 1.0.0 v stringr 1.5.1
## v lubridate 1.9.3 v tibble 3.2.1
## v readr 2.1.5
```

```
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag() masks stats::lag()
## x purrr::lift() masks caret::lift()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(rstatix)
```

```
##
## Attaching package: 'rstatix'
##
## The following object is masked _by_ '.GlobalEnv':
##
##     identify_outliers
##
## The following object is masked from 'package:stats':
##
##     filter
```

```
library(randomForest) #Random Forest
```

```
## Warning: package 'randomForest' was built under R version 4.4.1

## randomForest 4.7-1.2
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
##
## The following object is masked from 'package:dplyr':
##
##     combine
##
## The following object is masked from 'package:ggplot2':
##
##     margin
```

```
library(MASS) #Logistic Regression
```

```
##
## Attaching package: 'MASS'
##
## The following object is masked from 'package:rstatix':
##
##     select
##
## The following object is masked from 'package:dplyr':
##
##     select
```

```
load('data2/biomarker-clean.RData')
head(biomarker_clean)
```

```
## # A tibble: 6 x 1,319
##   group ados    CHIP CEBPB    NSE PIAS4 'IL-10 Ra' STAT3 IRF1 'c-Jun'
##   <chr> <dbl>  <dbl> <dbl>  <dbl> <dbl>    <dbl>  <dbl> <dbl>  <dbl>
## 1 ASD      8  0.335  0.520 -0.554  0.650    -0.358  0.305 -0.484  0.309
```

```
## 2 ASD      21 -0.0715  1.01   3      1.28      -0.133  1.13   0.253  0.408
## 3 ASD      12 -0.406  -0.531 -0.0592  1.13      0.554 -0.334  0.287 -0.845
## 4 ASD      20 -0.102  -0.251  1.47   0.0773     -0.705  0.893  2.61  -0.372
## 5 ASD      22 -0.395  -0.536  0.0410 -0.299     -0.830  0.899  1.01  -0.843
## 6 ASD      17 -0.126  1.27  -0.892  0.239     -0.344  0.216  0.211  0.221
## # i 1,309 more variables: 'Mcl-1' <dbl>, OAS1 <dbl>, 'c-Myc' <dbl>,
## #   SMAD3 <dbl>, SMAD2 <dbl>, 'IL-23' <dbl>, PDGFRA <dbl>, 'IL-12' <dbl>,
## #   STAT1 <dbl>, STAT6 <dbl>, LRRK2 <dbl>, Osteocalcin <dbl>, 'IL-5' <dbl>,
## #   GPDA <dbl>, IgA <dbl>, LPPL <dbl>, HEMK2 <dbl>, PDXK <dbl>, TLR4 <dbl>,
## #   REG4 <dbl>, 'HSP 27' <dbl>, 'YKL-40' <dbl>, 'Alpha enolase' <dbl>,
## #   'Apo L1' <dbl>, CD38 <dbl>, CD59 <dbl>, FABPL <dbl>, 'GDF-11' <dbl>,
## #   BTC <dbl>, 'HIF-1a' <dbl>, S100A6 <dbl>, SECTM1 <dbl>, RSP03 <dbl>, ...
```

```
set.seed(123)
```

```
# Split training (80%) and testing (20%) sets
trainIndex <- createDataPartition(biomarker_clean$group, p = 0.8, list = FALSE)
training_data <- biomarker_clean[trainIndex, ]
testing_data <- biomarker_clean[-trainIndex, ]
```

```
head(training_data)
```

Split Training and Testing

```
## # A tibble: 6 x 1,319
##   group ados   CHIP  CEBPB      NSE  PIAS4 'IL-10 Ra'  STAT3  IRF1 'c-Jun'
##   <chr> <dbl>  <dbl>  <dbl>    <dbl> <dbl>    <dbl>    <dbl> <dbl>  <dbl>
## 1 ASD      8  0.335  0.520 -0.554  0.650    -0.358  0.305 -0.484  0.309
## 2 ASD     22 -0.395 -0.536  0.0410 -0.299    -0.830  0.899  1.01  -0.843
## 3 ASD     17 -0.126  1.27  -0.892  0.239    -0.344  0.216  0.211  0.221
## 4 ASD     15  0.486  0.748 -1.09   0.462     0.570 -0.0682  1.01  1.21
## 5 ASD     10 -0.990 -1.10  0.231  -0.885    -0.151  0.0307 -0.0346 -0.891
## 6 ASD     22 -0.108  3      2.32   3      2.76  1.70  0.209  3
## # i 1,309 more variables: 'Mcl-1' <dbl>, OAS1 <dbl>, 'c-Myc' <dbl>,
## #   SMAD3 <dbl>, SMAD2 <dbl>, 'IL-23' <dbl>, PDGFRA <dbl>, 'IL-12' <dbl>,
## #   STAT1 <dbl>, STAT6 <dbl>, LRRK2 <dbl>, Osteocalcin <dbl>, 'IL-5' <dbl>,
## #   GPDA <dbl>, IgA <dbl>, LPPL <dbl>, HEMK2 <dbl>, PDXK <dbl>, TLR4 <dbl>,
## #   REG4 <dbl>, 'HSP 27' <dbl>, 'YKL-40' <dbl>, 'Alpha enolase' <dbl>,
## #   'Apo L1' <dbl>, CD38 <dbl>, CD59 <dbl>, FABPL <dbl>, 'GDF-11' <dbl>,
## #   BTC <dbl>, 'HIF-1a' <dbl>, S100A6 <dbl>, SECTM1 <dbl>, RSP03 <dbl>, ...
```

```
head(testing_data)
```

```
## # A tibble: 6 x 1,319
##   group ados   CHIP  CEBPB      NSE  PIAS4 'IL-10 Ra'  STAT3  IRF1 'c-Jun'
##   <chr> <dbl>  <dbl>  <dbl>    <dbl> <dbl>    <dbl>    <dbl> <dbl>  <dbl>
## 1 ASD     21 -0.0715  1.01   3      1.28     -0.133  1.13  0.253  0.408
## 2 ASD     12 -0.406  -0.531 -0.0592  1.13     0.554 -0.334  0.287 -0.845
```

```
## 3 ASD      20 -0.102 -0.251  1.47    0.0773    -0.705  0.893  2.61    -0.372
## 4 ASD      14 -0.378 -0.0790 -0.727  0.814    -0.811 -0.406 -0.791  -0.647
## 5 ASD      17  0.214  1.85    2.17    2.19     -0.102 -0.551 -0.293  1.80
## 6 ASD      13  1.35   -0.947 -1.28   -0.931    -0.443 -1.32  0.0259 -0.445
## # i 1,309 more variables: 'Mcl-1' <dbl>, OAS1 <dbl>, 'c-Myc' <dbl>,
## #   SMAD3 <dbl>, SMAD2 <dbl>, 'IL-23' <dbl>, PDGFRA <dbl>, 'IL-12' <dbl>,
## #   STAT1 <dbl>, STAT6 <dbl>, LRRK2 <dbl>, Osteocalcin <dbl>, 'IL-5' <dbl>,
## #   GPDA <dbl>, IgA <dbl>, LPPL <dbl>, HEMK2 <dbl>, PDXK <dbl>, TLR4 <dbl>,
## #   REG4 <dbl>, 'HSP 27' <dbl>, 'YKL-40' <dbl>, 'Alpha enolase' <dbl>,
## #   'Apo L1' <dbl>, CD38 <dbl>, CD59 <dbl>, FABPL <dbl>, 'GDF-11' <dbl>,
## #   BTC <dbl>, 'HIF-1a' <dbl>, S100A6 <dbl>, SECTM1 <dbl>, RSP03 <dbl>, ...
```

Apply T-test, Random Forest, and Logistic Regression (Using Top 20 Features)

```
library(dplyr)

# Ensure correct names and explicit call to dplyr::select
group_column <- training_data$group # Confirm column name is correct
protein_columns <- dplyr::select(training_data, -group, -ados) # Explicitly use dplyr

# T test for all protein
t_test_results <- sapply(protein_columns, function(protein) {
  t.test(protein ~ group_column)$p.value
})

# Convert result into frame
t_test_df <- data.frame(
  protein = names(t_test_results),
  p_value = t_test_results
)

# Extract top 20 proteins
top_proteins_ttest <- t_test_df %>%
  arrange(p_value) %>%
  slice(1:20) %>%
  pull(protein)

print(top_proteins_ttest)
```

T-test Selection

```
## [1] "PTN" "REL"
## [3] "MAPK2" "DERM"
## [5] "Calcineurin" "M2-PK"
## [7] "TFF3" "FSTL1"
## [9] "CXCL16, soluble" "MAPK14"
## [11] "Coagulation Factor IX" "IgD"
```

```
## [13] "MIA" "Fas, soluble"
## [15] "MMP-2" "IGFBP-4"
## [17] "ROR1" "Protein S"
## [19] "14-3-3 protein zeta/delta" "TGF-b R III"
```

```
library(randomForest)
library(dplyr)

# Alternative approach if select() conflicts remain
predictors <- training_data[, !names(training_data) %in% c("group", "ados")]
response <- factor(training_data$group)

set.seed(101422)

rf_model <- randomForest(x = predictors,
                          y = response,
                          ntree = 1000,
                          importance = TRUE)

print(rf_model$confusion)
```

Random Forest

```
## ASD TD class.error
## ASD 35 26 0.4262295
## TD 13 50 0.2063492
```

```
important_proteins_rf <- rf_model$importance %>%
  as_tibble() %>%
  mutate(protein = rownames(rf_model$importance)) %>%
  slice_max(MeanDecreaseGini, n = 20) %>% # top 20
  pull(protein)

print(important_proteins_rf)
```

```
## [1] "eIF-4H" "MAPK2" "PTN"
## [4] "CSK" "MAPK14" "M2-PK"
## [7] "DERM" "Lysozyme" "RELT"
## [10] "ILT-4" "CD27" "Nectin-like protein 2"
## [13] "Coagulation Factor IX" "Calcineurin" "Notch 1"
## [16] "IgD" "IGFBP-1" "SOST"
## [19] "GPVI" "MMP-2"
```



```

library(caret)
library(dplyr)

train_data_for_model <- training_data %>%
  dplyr::select(-ados) %>%
  dplyr::mutate(group = as.factor(group))

train_control <- trainControl(method = "none")

logistic_model <- train(group ~ .,
  data = train_data_for_model,
  method = "glm",
  family = "binomial",
  trControl = train_control)

```

Logistic Regression

```
## Warning: glm.fit: algorithm did not converge
```

```

logistic_coefficients <- coef(logistic_model$finalModel)

coeff_df <- as.data.frame(logistic_coefficients) %>%
  rownames_to_column(var = "protein") %>%
  filter(protein != "(Intercept)") %>%
  rename(coefficient = logistic_coefficients) %>%
  mutate(abs_coefficient = abs(coefficient))

top_20_proteins <- coeff_df %>%
  arrange(desc(abs_coefficient)) %>%
  slice(1:20) %>%
  pull(protein)

print(top_20_proteins)

```

```

## [1] "SMAD3" "HXK1"
## [3] "Myostatin" "'\\'HIF-1a\\'\'"
## [5] "CHKB" "ISLR2"
## [7] "CSH" "HHLA2"
## [9] "RNF43" "OAS1"
## [11] "'\\'TIMP-1\\'\'" "CD59"
## [13] "SMOC1" "'\\'14-3-3 protein beta/alpha\\'\'"
## [15] "SNP25" "LDLR"
## [17] "S100A4" "EFNB1"
## [19] "EFNB2" "STAT6"

```

```

all_top_proteins <- list(
  ttest = top_proteins_ttest,
  rf = important_proteins_rf,
  logreg = top_20_proteins
)

fuzzy_intersection_proteins <- all_top_proteins %>%
  unlist() %>%
  table() %>%
  .[, >= 2] %>%
  names()

print(fuzzy_intersection_proteins)

```

Fuzzy Interaction

```

## [1] "Calcineurin"          "Coagulation Factor IX" "DERM"
## [4] "IgD"                  "M2-PK"                 "MAPK14"
## [7] "MAPK2"                 "MMP-2"                 "PTN"
## [10] "RELT"

```

```

# Evaluate each set of selected proteins on testing data
evaluate_model_accuracy <- function(selected_proteins, model_name) {
  # Explicitly use dplyr::select() and tidyselect::all_of()
  predictors <- dplyr::select(testing_data, tidyselect::all_of(selected_proteins))
  response <- factor(testing_data$group)

  rf_test <- randomForest(x = predictors, y = response, ntree = 100)
  accuracy <- sum(diag(rf_test$confusion)) / sum(rf_test$confusion)

  cat("Accuracy of", model_name, "model on testing data:", accuracy, "\n")
}

# Evaluate individual methods
evaluate_model_accuracy(top_proteins_ttest, "T-test")

```

How are results affected by each modification?

```
## Accuracy of T-test model on testing data: 0.720524
```

```
evaluate_model_accuracy(important_proteins_rf, "Random Forest")
```

```
## Accuracy of Random Forest model on testing data: 0.7894737
```

```
#evaluate_model_accuracy(top_20_proteins, "Logistic Regression")
```

```
# Evaluate fuzzy intersection
```

```
evaluate_model_accuracy(fuzzy_intersection_proteins, "Fuzzy Intersection")
```

Accuracy of Fuzzy Intersection model on testing data: 0.8241758

Each method yields different predictive performance, with Random Forest and the Fuzzy Intersection performing similarly well. The fuzzy intersection can be a useful compromise when seeking a balance between different selection criteria, but in this case, it does not significantly outperform Random Forest alone.